

# thmm Package

July 1, 2017

**Type** Package

**Title** a tiny (but robust) discrete time univariate HMM library

**Version** 1.1

**Date** 2017-11-18

**Author** Alain Viari

**Maintainer** <alain.viari@inria.fr>

**Description** a small HMM library for discrete time univariate models  
with careful handling of underflow.

**License** GPL

**Depends** methods

**LazyData** TRUE

**RoxygenNote** 5.0.1

## R topics documented:

mstep.dnorm . . . . .	2
print.DTHmm . . . . .	2
thmm . . . . .	3
thmm.aic . . . . .	8
thmm.backward . . . . .	9
thmm.baumwelch . . . . .	10
thmm.bic . . . . .	12
thmm.bw.ctrl . . . . .	13
thmm.estep . . . . .	14
thmm.forward . . . . .	15
thmm.forward.backward . . . . .	16
thmm.init . . . . .	17
thmm.parameters . . . . .	19
thmm.simulate . . . . .	20
thmm.viterbi . . . . .	21
<b>Index</b>	<b>24</b>

---

mstep.dnorm	<i>M-step function for univariate normal distribution</i>
-------------	---

---

**Description**

used as the M-step function of [thmm.baumwelch](#) for univariate normal distribution.

**Usage**

```
mstep.dnorm(obs, cond, ctrl, mean, sd)
```

**Arguments**

obs	numerical vector of observations
cond	estimates of the conditional expectations as computed by <a href="#">thmm.estep</a>
ctrl	list of control parameters created by <a href="#">thmm.bw.ctrl</a>
mean	list of states means (may be missing)
sd	list of states sd (may be missing)

**Value**

list of updated parameters mean and sd

**See Also**

[thmm.baumwelch](#)

---

print.DTHmm	<i>print method for DTHmm</i>
-------------	-------------------------------

---

**Description**

print method for DTHmm

**Usage**

```
## S3 method for class 'DTHmm'
print(x, ...)
```

**Arguments**

x	a DTHmm object
...	further arguments passed to or from other methods

**Value**

invisible(x)

## Description

A tiny R package for discrete time hidden Markov models (DTHmm) with careful handling of underflow.

Strongly inspired by the 'HiddenMarkov' package from David Harte (<https://cran.r-project.org/web/packages/HiddenMarkov/index.html>) but numerically more stable (for forward/backward computations) and with different (hopefully simpler) API. (the 'HiddenMarkov' package has more features (including Markov modulated GLMs))

**Concepts:** A discrete time hidden Markov model is defined by:

- $m$  (discrete) states
- a probability density function density describing the (univariate or multivariate, continuous or discrete) distribution with different parameters for each state.
- a transition probability matrix (of size  $m \times m$ ) between states
- a vector (of length  $m$ ) of initial probability for each state

**important note:** for the sake of simplicity the remaining of this documentation as well as functions documentation, will assume that density is univariate. However the library can be used with multivariate densities as well (of course this comes at a price of a bit of complication). Please consult the **multivariate densities** section below.

Given a sequence of observed (univariate) random variable  $X$  noted  $\text{obs}=\text{obs}_1, \dots, \text{obs}_n$  this package solves the following problems:

1. Given the model parameters, compute the probability of  $\text{obs}$ . This problem is solved by the **forward and backward** algorithms.
2. Given the model parameters, find the most likely sequence of (hidden) states which could have generated  $\text{obs}$ . This problem is solved by the **Viterbi** algorithm.
3. Given  $\text{obs}$ , find the most likely set of density parameters and transition probabilities. This problem is solved by the **Baum-Welch** algorithm.

**forward algorithm:** is used to compute:

- $\alpha_{\{i,j\}} = \Pr\{X_1 = \text{obs}_1, \dots, X_i = \text{obs}_i, \text{state}_i = j \mid \text{hmm}\}$  that is the probability of seeing the partial sequence  $(\text{obs}_1, \dots, \text{obs}_i)$  and ending up in state  $j$  at time  $i$  for this hmm.

from which one can derive the **log-likelihood** of observation  $\text{obs}$  with this hmm:

- $\log(\Pr\{X_1 = \text{obs}_1, \dots, X_n = \text{obs}_n \mid \text{hmm}\})$

**backward algorithm:** is used to compute:

- $\beta_{\{i,j\}} = \Pr\{X_{\{i+1\}} = \text{obs}_{\{i+1\}}, \dots, X_n = \text{obs}_n \mid \text{state}_i = j, \text{hmm}\}$  that is the probability of the ending partial sequence  $(\text{obs}_{i+1}, \dots, \text{obs}_n)$  given that we started at state  $j$  at time  $i$ , for this hmm

**forward-backward algorithm:** putting everything together, one can further compute:

- $\gamma_{\{i,j\}} = \Pr\{\text{state}_i = j \mid \text{obs}, \text{hmm}\}$  that is the probability of being at time  $i$  in state  $j$  given this observation and this hmm

- $\rho_{\{i\}} = \Pr\{X_1 = \text{obs}_1, \dots, X_i = \text{obs}_i \mid \text{hmm}\}$  that is the probability of seeing the partial sequence  $(\text{obs}_1, \dots, \text{obs}_i)$  with this hmm

**Viterbi's algorithm:** The purpose of the Viterbi's algorithm is to determine the sequence of states  $(k_1^*, \dots, k_n^*)$  which maximises the joint distribution of the hidden states given the entire observed process. i.e:

- $(k_1^*, \dots, k_n^*) = \text{argmax}(\Pr(\text{state}_1=k_1, \dots, \text{state}_n=k_n, X_1=\text{obs}_1, \dots, X_n=\text{obs}_n \mid \text{hmm}))$

this also allows to compute:

- $\nu_{\{i+1, j\}} = \Pr(\text{state}_1=k_1^*, \dots, \text{state}_i=k_i^*, \text{state}_{i+1} = j, X_1=\text{obs}_1, \dots, X_i=\text{obs}_i \mid \text{hmm})$  where  $k_i^*$  is the optimal state at time  $i$ . (local decoding)
- $\text{viterbi\_score} = \Pr(\text{state}_1=k_1^*, \dots, \text{state}_n=k_n^*, \text{obs} \mid \text{hmm})$  the joint probability of optimal sequence of states and observation.
- $\text{probseq} = \Pr(\text{state}_1=k_1^*, \dots, \text{state}_n=k_n^* \mid \text{obs}, \text{hmm})$  the probability of the optimal sequence of state conditionally to the observation.

**Baum-Welch algorithm:** The purpose of the Baum-Welch algorithm is to estimate the HMM parameters (i.e. distribution parameters and transition probabilities) that best fit a given observation. This is a version of the EM algorithm that iteratively alternates two steps: the **E-step** and the **M-step**. The E-step is generic but the M-step depends upon each distribution. A M-step for Normal distribution is provided (as `mstep.dnorm`) and some others are given in examples. See [thmm.baumwelch](#) for more details (in particular how to control which parameters are actually adjusted).

**Implementation:** current implementation uses either pure R code or compiled C code. the default is to use (quicker) C code, R code has been kept for debugging purposes.

**Multivariate densities:** Instead of a univariate density function you may define a bi- or multivariate density function as well. The main difference is that instead of a numerical vector of observations, you should provide a list of tuples (each element in the tuple corresponds to one of the random variable). For the bivariate case, however, a trick is to use complex numbers instead of a list of doublets. This makes the definition of density and observation much simpler. Please see examples below as well as samples in the test directory.

## Examples

```
# -----
# Normal (gaussian) HMM with two states
# state 1 : mean=-1, sd=0.1
# state 2 : mean=+1, sd=0.1
# ergodic HMM with transition probability = 0.1
#
# full specification:
trans <- matrix(c(0.9, 0.1, 0.1, 0.9), nrow=2)
init <- c(0.5, 0.5)
hmm <- thmm.init(dnorm, trans, init, mean=c(-1,1), sd=c(0.1, 0.1))
#
# *same as previous* with simplified call:
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)

# simulate some sample
obs <- thmm.simulate(hmm, 100, .seed=0)

# run viterbi to retrieve states
vit <- thmm.viterbi(hmm, obs$values)
```

```

## Not run:
# compare actual and predicted states
plot(obs$values)
lines(thmm.parameters(hmm, "mean")[obs$states], col=3, lwd=5)
lines(thmm.parameters(hmm, "mean")[vit$states], col=2)
## End(Not run)

# -----
# poisson distribution
#
hmm <- thmm.init(dpois, 0.1, lambda=c(5, 10))
obs <- thmm.simulate(hmm, 100, .seed=0)
vit <- thmm.viterbi(hmm, obs$values)
## Not run:
plot(obs$values)
lines(thmm.parameters(hmm, "lambda")[obs$states], col=3, lwd=5)
lines(thmm.parameters(hmm, "lambda")[vit$states], col=2)
## End(Not run)

# -----
# mixture of gaussians
#
# we need to define the pdf
#
# note that this function is 'vectorized' on parameters (and x).
# the computation with log=TRUE will avoid underflow
#
dmixnorm <- function(x, mean=0, sd=1, prob=1, log=FALSE) {
  msp <- mapply(function(m, s, p) list(mean=m, sd=s, prob=p), mean, sd, prob, SIMPLIFY=FALSE)
  res <- if (log) {
    lapply(x, function(x) sapply(msp, function(p) {
      pp <- base::log(p$prob) + dnorm(x, mean=p$mean, sd=p$sd, log=TRUE)
      mp <- max(pp)
      mp + base::log(sum(exp(pp-mp)))
    }))
  } else {
    lapply(x, function(x) sapply(msp, function(p)
      sum(p$prob*dnorm(x, mean=p$mean, sd=p$sd, log=FALSE))))
  }
  unlist(res)
}
#
# and random generator
# again 'vectorized' on parameters (and x).
rmixnorm <- function(n, mean=0, sd=1, prob=1) {
  .cycle <- function(x, n) head(rep(x, n), n)
  msp <- mapply(function(m, s, p) list(mean=m, sd=s, prob=p), mean, sd, prob, SIMPLIFY=FALSE)
  unlist(lapply(seq_len(n), function(i) sapply(msp, function(p) {
    r <- mapply(function(m, s) rnorm(1, m, s), p$mean, p$sd)
    sample(r, size=1, prob=.cycle(p$prob, length(r)))
  })))
}

# hmm with 2 states
# state1: mixture of 2 gaussians with mean=-1,1 sd=0.1,0.1 and prob=0.1,0.9
# state2: single gaussian mean=0 sd=0.2

```

```

hmm <- thmm.init(dmixnorm, 0.2, mean=list(c(-1,1), 0), sd=list(0.1, 0.2),
                prob=list(c(0.5,0.5), 1))
obs <- thmm.simulate(hmm, n=100, .seed=0)
vit <- thmm.viterbi(hmm, obs$values)
## Not run:
plot(obs$values)
lines((2-obs$states), col=3)
y <- thmm.parameters(hmm, "mean")[vit$states]
lines(sapply(y, function(x) x[2%%length(x)+1]), col=2)
## End(Not run)

# -----
# discrete distribution : the dishonest casino
#
# the pdf (with p6 parameter = probability of drawing a six)
dice <- function(x, p6=1/6, log=FALSE) {
  r <- unlist(lapply(x, function(x) {
    unlist(lapply(p6, function(p6) {
      if (is.na(x)) NA
      else if (x == 6) p6
      else if (x %in% 1:5) (1-p6)/5
      else 0
    }) })), recursive=FALSE)
  if (log) r <- base::log(r)
  r
}
# and the random generator (d->r)ice
rice <- function(n, p6=1/6) {
  sample(1:6, n, prob=c(rep((1-p6)/5, 5), p6), replace=TRUE)
}

hmm <- thmm.init(dice, 0.1, p6=c(1/6, 3/6)) # the second dice is loaded
obs <- thmm.simulate(hmm, 100, .seed=0)
vit <- thmm.viterbi(hmm, obs$values)
## Not run:
plot(obs$values)
lines((obs$states-1)*5+1, col=3) # truth
lines((vit$states-1)*5+1, col=2) #predicted
## End(Not run)

# -----
# univariate normal : Baum-Welch
#
# reference hmm
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)
obs <- thmm.simulate(hmm, n=1000, .seed=0)

# this one converge to the correct solution
hmm0 <- thmm.init(dnorm, 0.1, mean=c(0,5))
bw <- thmm.baumwelch(hmm0, obs$values)

# but not that one
hmm0 <- thmm.init(dnorm, 0.1, mean=c(0,10))
bw <- thmm.baumwelch(hmm0, obs$values)

# so adding constraints: no update of trans nor init
hmm0 <- thmm.init(dnorm, 0.1, mean=c(0,10))

```

```

ctrl <- thmm.bw.ctrl(hmm0, do.trans=FALSE, do.init=FALSE)
bw <- thmm.baumwelch(hmm0, obs$values, ctrl)

# -----
# bivariate normal
# implemented using complex numbers
# see test.multivar.norm.r for a more general multivariate version
#

# helper
.c <- function(x,y=x) complex(real=x, imaginary=y)

#
# pdf bivariate normal (with 0 covariance)
#
dxynorm <- function(x, mean=.c(0), sd=.c(1), log=FALSE) {
  re <- dnorm(Re(x), mean=Re(mean), sd=Re(sd), log=log)
  im <- dnorm(Im(x), mean=Im(mean), sd=Im(sd), log=log)
  if (log) re+im else re*im
}

#
# random generation function
#
rxynorm <- function(n, mean=.c(0), sd=.c(1)) {
  re <- rnorm(n, mean=Re(mean), sd=Re(sd))
  im <- rnorm(n, mean=Im(mean), sd=Im(sd))
  .c(re, im)
}

hmm <- thmm.init(dxynorm, 0.1,
                 mean=c(.c(10, 20), .c(50, 100)),
                 sd=c(.c(5,10),.c(25, 40)))

obs <- thmm.simulate(hmm, 100, with.values=TRUE)
vit <- thmm.viterbi(hmm, obs$values)
lvl <- thmm.parameters(hmm, "mean")

## Not run:
par(mfrow=c(2,1))
plot(Re(obs$values))
lines(Re(lvl[obs$states]), col=3, lwd=2)
lines(Re(lvl[vit$states]), col=2, lty=2, lwd=2)
plot(Im(obs$values))
lines(Im(lvl[obs$states]), col=3, lwd=2)
lines(Im(lvl[vit$states]), col=2, lty=2, lwd=2)
par(mfrow=c(1,1))
## End(Not run)

# -----
# bivariate normal + poisson
#

# helper : take ith element of tuple
.i <- function(x, i) sapply(x, function(x) x[i])

```

```

# pdf : bivariate normal + poisson
#

dnorpois <- function(x, mean=0, sd=1, lambda=10, log=FALSE) {
  .prod <- if (log) sum else prod
  unlist(mapply(function(m, s, l) {
    xn <- dnorm(.i(x,1), m, s, log=log)
    xp <- dpois(.i(x,2), l, log=log)
    mapply(.prod, xn, xp)
  },
    mean, sd, lambda, SIMPLIFY=FALSE), use.names=FALSE)
}

# random number generator
#

rnormpois <- function(n, mean=0, sd=1, lambda=10) {
  mapply(c, rnorm(n, mean, sd), rpois(n, lambda), SIMPLIFY=FALSE)
}

hmm <- thmm.init(dnorpois, 0.1,
  mean=c(10, 20, 30), sd=c(1, 2, 3),
  lambda=c(5, 10, 15))

obs <- thmm.simulate(hmm, 100, with.values=TRUE)

vit <- thmm.viterbi(hmm, obs$values)

lvl.mean <- thmm.parameters(hmm, "mean")
lvl.lamb <- thmm.parameters(hmm, "lambda")

## Not run:
par(mfrow=c(2,1))
plot(.i(obs$values,1), main="normal")
lines(lvl.mean[obs$states], col=3, lwd=2)
lines(lvl.mean[vit$states], col=2, lty=2, lwd=2)
plot(.i(obs$values,2), main="poisson")
lines(lvl.lamb[obs$states], col=3, lwd=2)
lines(lvl.lamb[vit$states], col=2, lty=2, lwd=2)
par(mfrow=c(1,1))
## End(Not run)

```

thmm.aic

*compute AIC*

## Description

compute Akaike's Information Criterion (AIC) of hmm and observation.

## Usage

```
thmm.aic(hmm, obs, np, nt = nrow(hmm$trans) * (nrow(hmm$trans) - 1),
  ni = nrow(hmm$trans) - 1, k = 2, llk = NA, .useC = TRUE)
```



**Arguments**

hmm	DTHmm model (from <a href="#">thmm.init</a> )
obs	numerical vector of observations (may be NULL if llk is provided)
np	number of free parameters of hmm density function (e.g. 2 for normal, 1 for poisson).
nt	number of free parameters for transition matrix.
ni	number of free parameters for initial probabilities.
k	the penalty parameter. the default k=2 is the classical AIC.
llk	numeric, log-likelihood(obs   hmm). if NA then llk will be computed using <a href="#">thmm.forward</a> .
.useC	logical if TRUE (default) use C code else use R code for <a href="#">thmm.forward</a>

**Details**

$AIC = -2 * llk + (nt + ni + M * np) * k$   
 with M=number of states

**Value**

numeric AIC

**Examples**

```
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)
obs <- thmm.simulate(hmm, 100, .seed=0)$values
thmm.aic(hmm, obs, np=2, nt=1, ni=0)
```

---

thmm.backward	<i>(log) backward probabilities</i>
---------------	-------------------------------------

---

**Description**

compute backward probabilities of observations according to a discrete time hidden markov model (DTHmm).

**Usage**

```
thmm.backward(hmm, obs, .useC = TRUE)
```

**Arguments**

hmm	DTHmm model (from <a href="#">thmm.init</a> )
obs	numerical vector of observations
.useC	logical if TRUE (default) use C code else use R code

**Value**

list with one components: logbeta (for symmetry with [thmm.forward](#))

- logbeta : the (log of) backward probabilities (densities) matrix:  
 $\text{beta}_{\{i,j\}} = \Pr\{X_{\{i+1\}} = \text{obs}_{\{i+1\}}, \dots, X_n = \text{obs}_n \mid \text{state}_i = j, \text{hmm}\}$   
 that is the probability (density) of the ending partial sequence ( $\text{obs}_{i+1}, \dots, \text{obs}_n$ ) given that we started at state  $j$  at time  $i$ , for this hmm.

**Note**

the code takes care of rescaling values during calculation to avoid underflow problems and is quite robust in practice.

**See Also**

[thmm.forward](#), [thmm.forward.backward](#)

**Examples**

```
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)
obs <- c(rnorm(100, -1, 0.1), rnorm(100, 1, 0.1))
bwd <- thmm.backward(hmm, obs)
#
# test numerical stability
#
dif1 <- bwd$logbeta[-c(200,99),1] - bwd$logbeta[-c(200,99),2]
obs[100] <- 1000000 # this may creates overflow
bwd <- thmm.backward(hmm, obs)
dif2 <- bwd$logbeta[-c(200,99),1] - bwd$logbeta[-c(200,99),2]
max(abs((dif1-dif2)/dif1))
## Not run:
plot(dif1)
points(dif2, col=2, cex=0.1)
## End(Not run)
```

---

thmm.baumwelch

*Baum-Welch algorithm*


---

**Description**

HMM parameters estimation using the Baum-Welch algorithm

**Usage**

```
thmm.baumwelch(hmm0, obs, ctrl = thmm.bw.ctrl(hmm0), .mstep, .useC = TRUE)
```

**Arguments**

hmm0	a starting DTHmm (see <a href="#">thmm.init</a> )
obs	numerical vector of observations
ctrl	a list of control settings obtained by <a href="#">thmm.bw.ctrl</a>
.mstep	function name of the mstep function (see details)
.useC	logical if TRUE (default) use C code else use R code

## Details

the Baum-Welch EM algorithm iteratively alternates two steps: the **E-step** and the **M-step**. The **E-step** (expectation) is distribution independent and is performed by the [thmm.estep](#) function. The **M-step** (maximisation) is distribution dependent and you should provide a distribution specific function to perform it. The name of this function is `mstep.<density>` (e.g. `mstep.dnorm` for a gaussian distribution) unless you provide another name as the `.mstep` parameter. Default function is provided for `dnorm` (may be extended in the future). Have a look at [mstep.dnorm](#) for an example of function format.

you may control which hmm parameters are actually optimized with `ctrl`. see [thmm.bw.ctrl](#).

## Value

a named list of three elements:

- `hmm` : optimized hmm
- `info` : information about the iterative process
- `loglike` : final log-likelihood

## Examples

```
# reference hmm
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)
obs <- thmm.simulate(hmm, n=1000, .seed=0)

# this one converge to the correct solution
hmm0 <- thmm.init(dnorm, 0.1, mean=c(0,5))
bw <- thmm.baumwelch(hmm0, obs$values)

# but not that one
hmm0 <- thmm.init(dnorm, 0.1, mean=c(0,10))
bw <- thmm.baumwelch(hmm0, obs$values)

# so adding constraints: no update of trans nor init
hmm0 <- thmm.init(dnorm, 0.1, mean=c(0,10))
ctrl <- thmm.bw.ctrl(hmm0, do.trans=FALSE, do.init=FALSE, verbose=TRUE)
bw <- thmm.baumwelch(hmm0, obs$values, ctrl)

# that one do not converge either to proper solution
hmm0 <- thmm.init(dnorm, 0.1, mean=c(0,0), sd=0.5)
bw <- thmm.baumwelch(hmm0, obs$values)

# adding constraints: mean bounds
hmm0 <- thmm.init(dnorm, 0.1, mean=c(0,0), sd=0.5)
ctrl <- thmm.bw.ctrl(hmm0, lower.mean=c(-2,0), upper.mean=c(0,2), verbose=TRUE)
bw <- thmm.baumwelch(hmm0, obs$values, ctrl)

# adding constraint: as as previous but more general form
hmm0 <- thmm.init(dnorm, 0.1, mean=c(0,0), sd=0.5)
myconst <- function(params) {
  params$mean <- pmax(pmin(params$mean, c(0,2)), c(-2,0))
  params
}
ctrl <- thmm.bw.ctrl(hmm0, constraint=myconst, verbose=TRUE)
bw <- thmm.baumwelch(hmm0, obs$values, ctrl)
```

---

thmm.bic	<i>compute BIC</i>
----------	--------------------

---

### Description

compute Bayesian Information Criterion (BIC) of hmm and observation.

### Usage

```
thmm.bic(hmm, obs, np, nt = nrow(hmm$trans) * (nrow(hmm$trans) - 1),
  ni = nrow(hmm$trans) - 1, llk = NA, .useC = TRUE)
```

### Arguments

hmm	DTHmm model (from <a href="#">thmm.init</a> )
obs	numerical vector of observations
np	number of free parameters of hmm density function (e.g. 2 for normal, 1 for poisson).
nt	number of free parameters for transition matrix.
ni	number of free parameters for initial probabilities.
llk	numeric, log-likelihood(obs   hmm). if NA then llk will be computed using <a href="#">thmm.forward</a> .
.useC	logical if TRUE (default) use C code else use R code for <a href="#">thmm.forward</a>

### Details

$BIC = -2 * llk + (nt + ni + M * np) * \log(N)$   
 with M=number of states; N=length(obs)

### Value

numeric BIC

### Examples

```
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)
obs <- thmm.simulate(hmm, 100, .seed=0)$values
thmm.bic(hmm, obs, np=2, nt=1, ni=0)
```

thmm.bw.ctrl

*control parameters for Baum-Welch algorithm***Description**

make a list of parameters to control [thmm.baumwelch](#) execution.

**Usage**

```
thmm.bw.ctrl(hmm, maxiter = NA, tol = 1e-05, verbose = FALSE,
  converge = expression(difLL <= ctrl$tol), do.trans = TRUE,
  do.init = TRUE, constraint = NULL, ...)
```

**Arguments**

hmm	DTHmm model (from <a href="#">thmm.init</a> )
maxiter	integer max number of iteration (NA will set to hmm dependent value)
tol	numeric convergence criterion (delta_log-likelihood)
verbose	logical be verbose
converge	expression of the convergence criterion
do.trans	logical do optimize transition matrix
do.init	logical do optimize initial probabilities
constraint	if not NULL, should contain a function receiving a list of optimized parameters (at each M-step) and should return a list (with the same format) of constrained parameters values (see examples).
...	any parameter passed to <density>.mstep function (see <a href="#">thmm.baumwelch</a> ). in particular: <ul style="list-style-type: none"> <li>parameters of the form do.&lt;parameter&gt; are interpreted as logical controlling whether or not we should optimize the &lt;parameter&gt;. note that to prevent update you should explicitly set do.&lt;parameter&gt; to FALSE. if do.&lt;parameter&gt; is absent (or set to TRUE) then &lt;parameter&gt; <b>will be</b> updated.</li> <li>parameters of the form upper lower.&lt;parameter&gt; are interpreted as numerical upper/lower bounds. if absent then no bounds (i.e. [-Inf, +Inf]) are assumed.</li> </ul>

**Value**

named list of control parameters for [thmm.baumwelch](#)

**Note**

constraint provides a more general form of control than do.<parameter>, upper|lower.<parameter> but may be less efficient since it is called right after the mstep function.

(for developers) do.<parameter>, upper|lower.<parameter> should be implemented for every mstep function. see [mstep.dnorm](#) as an exemple. constraint is called by [thmm.baumwelch](#) directly.

**See Also**

[thmm.baumwelch](#)

---

thmm.estep	<i>E-Step of EM Baum-Welch algorithm</i>
------------	--

---

**Description**

performs the expectation step of the EM algorithm for a discrete time HMM process. this function is called by the [thmm.baumwelch](#).

**Usage**

```
thmm.estep(hmm, obs, .useC = TRUE)
```

**Arguments**

hmm	DTHmm model (from <a href="#">thmm.init</a> )
obs	numerical vector of observations
.useC	logical if TRUE (default) use C code else use R code

**Details**

u is defined as:  
 $u_{i,j} = \Pr(\text{state}_i = j \mid \text{obs}, \text{hmm})$   
 that is the same as gamma in [thmm.forward.backward](#)

v is defined as:  
 $v_{\{i,j,k\}} = \Pr(\text{state}_{i-1} = j, \text{state}_i = k \mid \text{obs}, \text{hmm})$   
 also known as chsi

**Value**

list with three components:

- u  $\text{length}(\text{obs}) \times \text{number\_states}$  matrix containing estimates of the conditional expectations (see details)
- v  $\text{length}(\text{obs}) \times \text{number\_states} \times \text{number\_states}$  matrix containing estimates of the conditional expectations (see details)
- loglike log-likelihood

**See Also**

[thmm.baumwelch](#)

---

thmm.forward	(log) forward probabilities
--------------	-----------------------------

---

## Description

compute forward probabilities of observations according to a discrete time hidden markov model (DTHmm).

## Usage

```
thmm.forward(hmm, obs, .useC = TRUE)
```

## Arguments

hmm	DTHmm model (from <a href="#">thmm.init</a> )
obs	numerical vector of observations
.useC	logical if TRUE (default) use C code else use R code

## Value

list with two components: logalpha and loglike

- logalpha : the (log of) forward probabilities (densities) matrix:  
 $\alpha_{i,j} = \Pr\{X_1 = \text{obs}_1, \dots, X_i = \text{obs}_i, \text{state}_i = j \mid \text{hmm}\}$  that is the probability (density) of seeing the partial sequence (obs\_1, ..., obs\_i) and ending up in state j at time i for this hmm.
- loglike : the log-likelihood of this observation with this hmm:  
 $\log(\Pr\{X_1 = \text{obs}_1, \dots, X_n = \text{obs}_n \mid \text{hmm}\}) = \max_j (\log \alpha_{n,j})$

## Note

the code takes care of rescaling values during calculation to avoid underflow problems and is quite robust in practice.

since the computation uses density probabilities it is perfectly valid to get a positive log-likelihood.  
 for multivariate HMM, obs should be a list of (numerical) tuples. (see [thmm](#) examples).

## See Also

[thmm.backward](#), [thmm.forward.backward](#)

## Examples

```
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)
obs <- c(rnorm(100, -1, 0.1), rnorm(100, 1, 0.1))
fwd <- thmm.forward(hmm, obs)
#
# test numerical stability
#
dif1 <- fwd$logalpha[-100,1] - fwd$logalpha[-100,2]
obs[100] <- 1000000 # this may creates overflow
fwd <- thmm.forward(hmm, obs)
```

```
dif2 <- fwd$logalpha[-100,1] - fwd$logalpha[-100,2]
max(abs((dif1-dif2)/dif1))
## Not run:
plot(dif1)
points(dif2, col=2, cex=0.1)
## End(Not run)
```

---

thmm.forward.backward (log) forward/backward probabilities

---

## Description

compute forward/backward probabilities of observations according to a discrete time hidden markov model (DTHmm).

## Usage

```
thmm.forward.backward(hmm, obs, .useC = TRUE)
```

## Arguments

hmm	DTHmm model (from <a href="#">thmm.init</a> )
obs	numerical vector of observations
.useC	logical if TRUE (default) use C code else use R code

## Value

list with five components: logalpha, logbeta, gamma, rho, loglike

- **logalpha** : the (log of) forward probabilities (densities) matrix:  
 $\alpha_{\{i,j\}} = \Pr\{X_1 = \text{obs}_1, \dots, X_i = \text{obs}_i, \text{state}_i = j \mid \text{hmm}\}$  that is the probability (density) of seeing the partial sequence ( $\text{obs}_1, \dots, \text{obs}_i$ ) and ending up in state  $j$  at time  $i$  for this hmm.
- **logbeta** : the (log of) backward probabilities (densities) matrix:  
 $\beta_{\{i,j\}} = \Pr\{X_{\{i+1\}} = \text{obs}_{\{i+1\}}, \dots, X_n = \text{obs}_n \mid \text{state}_i = j, \text{hmm}\}$  that is the probability (density) of the ending partial sequence ( $\text{obs}_{i+1}, \dots, \text{obs}_n$ ) given that we started at state  $j$  at time  $i$ , for this hmm.
- **gamma** : the probabilities matrix:  
 $\gamma_{\{i,j\}} = \Pr\{\text{state}_i = j \mid \text{obs}, \text{hmm}\}$  that is the probability of being at time  $i$  in state  $j$  given this observation and this hmm.
- **logrho** : the (log of) probabilities (densities) vector:  
 $\rho_{\{i\}} = \Pr\{X_1 = \text{obs}_1, \dots, X_i = \text{obs}_i \mid \text{hmm}\}$  that is the probability (density) of seeing the partial sequence ( $\text{obs}_1, \dots, \text{obs}_i$ ) for this hmm.
- **loglike** : the log-likelihood of this observation with this hmm:  
 $\log(\Pr\{X_1 = \text{obs}_1, \dots, X_n = \text{obs}_n \mid \text{hmm}\})$  (see [thmm.forward](#))

## Note

the code takes care of rescaling values during calculation to avoid underflow problems and is quite robust in practice.

by definition `loglike = logrho(length(obs))` (although it is not computed this way).



**See Also**

[thmm.forward](#), [thmm.backward](#)

**Examples**

```
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)
obs <- c(rnorm(100, -1, 0.1), rnorm(100, 1, 0.1))
fb <- thmm.forward.backward(hmm, obs)
#
# test numerical stability
#
obs[101] <- 1000000 # this may creates overflow
fb <- thmm.forward.backward(hmm, obs)
## Not run:
plot(fb$gamma[,2])
## End(Not run)
```

---

thmm.init	<i>initialize HMM</i>
-----------	-----------------------

---

**Description**

initialize a discrete time hidden markov model (DTHmm)

**Usage**

```
thmm.init(density, trans, init = NULL, ...)
```

**Arguments**

density	a probability density function (pdf) such as <a href="#">dnorm</a> (see notes)
trans	the transition matrix given either as a matrix, a vector or a constant (see details)
init	vector of probabilities of the initial states (see details)
...	parameters of density (see examples)

**Details**

the number of states (m) is determined from the ... arguments (see notes) and should be consistent with the value passed to trans and init if provided.

if trans is a constant then the transition matrix is computed as a mxm matrix with **off-diagonal** elements trans.

if trans is a vector (it should then be of length m) then the transition matrix is computed as a mxm matrix with **diagonal** elements trans.

if init is NULL then init is computed as a vector of uniform probabilities (1/m).

if init is not of length m then it is recycled/shortened to proper length m.

**Value**

a DTHmm object describing the DTHmm

**Note**

the pdf function density has the form:

`density(x, ..., log=TRUE|FALSE)`

see [dnorm](#), [dpois](#), [dbeta](#), etc. as examples but you may write your own as well, with the following specifications:

- density parameters must have default values if you don't specify them in ...
- density must accept vectors (including of NA's) as parameters (possibly recycling them) (each parameter corresponds to each state) and return a vector of length m (possibly of NA's) when called with a scalar x (including NA). this is used to determine the actual number of states m.
- (not mandatory) It may accept a vector (including of NA's) as x argument (with individual states parameters) and return a vector of size `length(x)`. this is used by the C code only, so if this condition is not satisfied the C code will workaround but will run a bit slower (albeit still quicker than R code).

You may pass the function symbol or function name (character) as density argument (but not a direct lambda expression nor anything fancy).

For multivariate HMM, the density function should accept a list of (numeric) tuples as x argument (or anything equivalent like a list of complex numbers for bivariate). The parameters can be anything relevant to your multivariate density distribution but may be received as lists (one element for each state) instead of vectors. See examples in [thmm](#) for details.

**Examples**

```
# -----
# Normal (gaussian) HMM with two states
# full specification
trans <- matrix(c(0.9, 0.1, 0.1, 0.9), nrow=2)
init <- c(0.5, 0.5)
hmm <- thmm.init(dnorm, trans, init, mean=c(-1,1), sd=c(0.1, 0.1))

# -----
# *same as previous* with simplified call
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)

# -----
# three states HMM with Poisson distribution
hmm <- thmm.init(dpois, 0.1, lambda=1:3)

# -----
# two states HMM with Gamma distribution
hmm <- thmm.init(dgamma, 0.1, shape=c(3, 10), rate=c(1, 2))

# -----
# mixture of gaussians
# state1: mixture of 2 gaussians with mean=-1,1 sd=0.1,0.1 and prob=0.1,0.9
# state2: gaussian mean=0 sd=0.2
#
# we define first \code{dmixnorm} as the pdf of a mixture of gaussians
# note that this function is 'vectorized' on parameters (and x either).
# the computation with log=TRUE will avoid underflow
#
```

```

dmixnorm <- function(x, mean=0, sd=1, prob=1, log=FALSE) {
  msp <- mapply(function(m, s, p) list(mean=m, sd=s, prob=p), mean, sd, prob, SIMPLIFY=FALSE)
  res <- if (log) {
    lapply(x, function(x) sapply(msp, function(p) {
      pp <- base::log(p$prob) + dnorm(x, mean=p$mean, sd=p$sd, log=TRUE)
      mp <- max(pp)
      mp + base::log(sum(exp(pp-mp)))
    }))
  } else {
    lapply(x, function(x) sapply(msp, function(p)
      sum(p$prob*dnorm(x, mean=p$mean, sd=p$sd, log=FALSE))))
  }
  unlist(res)
}
#
# then the hmm
hmm <- thmm.init(dmixnorm, 0.1, mean=list(c(-1,1), 0), sd=list(0.1, 0.2),
  prob=list(c(0.5,0.5), 1))

# -----
# discrete distribution : binomial
#
# dbinom does not have default parameters
# so we need to specify all of them
hmm <- thmm.init(dbinom, 0.1, size=c(50, 100), prob=c(0.5, 0.1))

# -----
# discrete distribution : the dishonest casino
#
# we first define the pdf
#
dice <- function(x, p6=1/6, log=FALSE) {
  r <- unlist(lapply(x, function(x) {
    unlist(lapply(p6, function(p6) {
      if (is.na(x)) NA
      else if (x == 6) p6
      else if (x %in% 1:5) (1-p6)/5
      else 0
    }))) }, recursive=FALSE)
  if (log) r <- base::log(r)
  r
}
# and the random generator (d->r)ice (see \link{thmm.simulate})
rice <- function(n, p6=1/6) {
  sample(1:6, n, prob=c(rep((1-p6)/5, 5), p6), replace=TRUE)
}

hmm <- thmm.init(dice, 0.3, p6=c(1/6, 3/6))
obs <- thmm.simulate(hmm, 100)
## Not run:
plot(obs$values)

## End(Not run)

```

**Description**

get hmm global and individual states parameters

**Usage**

```
thmm.parameters(hmm, what = NULL)
```

**Arguments**

hmm	DTHmm model (from <a href="#">thmm.init</a> )
what	character parameter name to retrieve (all parameters if NULL (default))

**Value**

if what==NULL return a list of 5 components: nstates, density, args, init, trans

- nstates : integer number of states
- density : character string name of probability density function
- args : list (of length nstates) of named lists, each element gives the density function parameters for state i. (see note).
- init : numerical vector of initial probabilities
- trans : numerical transition matrix

if what!=NULL then what should be either one of the above names or the name of one of the density function parameters. the function then return only this information.

**Examples**

```
# univariate gaussian hmm
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=1)
ns <- thmm.parameters(hmm, "nstates")
means <- thmm.parameters(hmm, "mean")
# viterbi decode
obs <- rnorm(100)
vit <- thmm.viterbi(hmm, obs)
# as values
means[vit$states]
```

---

thmm.simulate

*simulate hidden markov process*


---

**Description**

generate states and observations vectors following a discrete time HMM

**Usage**

```
thmm.simulate(hmm, npts, with.values = TRUE, log = FALSE, .random = NULL,
  .seed = NA)
```

**Arguments**

<code>hmm</code>	DTHmm model (from <a href="#">thmm.init</a> )
<code>npts</code>	number of points to simulate
<code>with.values</code>	logical also compute observed values (see details)
<code>log</code>	(if <code>with.values=TRUE</code> ) compute log of values
<code>.random</code>	name of random generation function (see details).
<code>.seed</code>	if not NA set random seed before generation

**Details**

if `with.values=TRUE` then the random generation function used to produce values from states is '(d->r)ensity' i.e. the name of density probability function with first letter (usually a 'd') replaced by a 'r'. (e.g. if `density=dnorm` then the random function is `rnorm`) this can be overridden by using the `.random` parameter.

This function has the following form:

`random(n, ...)` where ... are the same parameters as `density`.

In contrast to `density`, the generation function does not need to be vectorized on parameters nor `n` (actually it is always called with `n=1` and individual state parameters).

**Value**

if `with.values=FALSE` integer vector of simulated states. if `with.values=TRUE` a list of two vectors states and values.

**Note**

neither states nor values are deterministic, you should call [set.seed](#) to produce reproducible vectors.

for multivariate HMM, the result is a list of (numeric) tuples (one tuple element per each random variable). See examples in [thmm](#) for details.

**Examples**

```
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)
set.seed(0)
thmm.simulate(hmm, 10, log=TRUE)
```

---

thmm.viterbi

*Viterbi algorithm*


---

**Description**

This function calculates the optimal hidden states sequence using Viterbi's algorithm.

**Usage**

```
thmm.viterbi(hmm, obs, with.probseq = FALSE, .useC = TRUE)
```

**Arguments**

<code>hmm</code>	DTHmm model (from <a href="#">thmm.init</a> )
<code>obs</code>	numerical vector of observations
<code>with.probseq</code>	should we compute logprobseq and loglike (see value)
<code>.useC</code>	logical if TRUE (default) use C code else use R code

**Details**

the purpose of the Viterbi's algorithm is to determine the sequence of states ( $k_1^*, \dots, k_n^*$ ) which maximises the joint distribution of the hidden states given the entire observed process. i.e:

$$(k_1^*, \dots, k_n^*) = \operatorname{argmax}(\Pr(\text{state}_1=k_1, \dots, \text{state}_n=k_n, X_1=\text{obs}_1, \dots, X_n=\text{obs}_n \mid \text{hmm}))$$
**Value**

list with five components: `states`, `lognu`, `logviterbi` `loglike` and `logprobseq`

- `states` : vector of optimal states ( $k_1^*, \dots, k_n^*$ )
- `lognu` : the (log of) joint probabilities (densities) matrix:  
 $\text{nu}_{\{i+1,j\}} = \Pr(\text{state}_1=k_1^*, \dots, \text{state}_i=k_i^*, \text{state}_{i+1}=j, X_1=\text{obs}_1, \dots, X_i=\text{obs}_i \mid \text{hmm})$   
 where  $k_i^*$  is the optimal state at time  $i$ .
- `logviterbi` : log of Viterbi's score, i.e. the log of joint probability (density) of the optimal sequence of states and observations with this hmm  
 $\log(\Pr(\text{state}_1=k_1^*, \dots, \text{state}_n=k_n^*, \text{obs} \mid \text{hmm}))$   
 $= \max_j(\log \text{nu}_{\{n,j\}})$
- `loglike` : the log-likelihood of this observation with this hmm:  
 $\log(\Pr\{X_1 = \text{obs}_1, \dots, X_n = \text{obs}_n \mid \text{hmm}\}) = \log(\Pr(\text{obs} \mid \text{hmm}))$  (see [thmm.forward](#))  
 this value is NA if `with.probseq == FALSE`
- `logprobseq` : log of probability (density) of the optimal sequence of states conditionally to the observations. i.e  
 $\log(\Pr(\text{state}_1=k_1^*, \dots, \text{state}_n=k_n^* \mid \text{obs}, \text{hmm}))$   
 $= \log(\Pr(\text{state}_1=k_1^*, \dots, \text{state}_n=k_n^*, \text{obs} \mid \text{hmm}) / \Pr(\text{obs} \mid \text{hmm}))$   
 $= \text{logviterbi} - \text{loglike}$

this value is NA if `with.probseq == FALSE`

**Note**

the code takes care of rescaling values during calculation to avoid underflow problems and is quite robust in practice.

for multivariate HMM, `obs` should be a list of (numerical) tuples. (see [thmm](#) examples).

**See Also**

[thmm.forward](#)

**Examples**

```
hmm <- thmm.init(dnorm, 0.1, mean=c(-1,1), sd=0.1)
obs <- c(rnorm(100, -1, 0.1), rnorm(100, 1, 0.1))
vit <- thmm.viterbi(hmm, obs, with.probseq=TRUE)
table(vit$states)
#
```

```
# test numerical stability
#
obs[101] <- 1000000 # this may creates overflow
vit <- thmm.viterbi(hmm, obs, with.probsub=TRUE)
table(vit$states)
## Not run:
plot(vit$states)
## End(Not run)
```

# Index

dbeta, [18](#)  
dnorm, [17](#), [18](#)  
dpois, [18](#)  
  
mstep.dnorm, [2](#), [11](#), [13](#)  
  
print.DTHmm, [2](#)  
  
set.seed, [21](#)  
  
thmm, [3](#), [15](#), [18](#), [21](#), [22](#)  
thmm.aic, [8](#)  
thmm.backward, [9](#), [15](#), [17](#)  
thmm.baumwelch, [2](#), [4](#), [10](#), [13](#), [14](#)  
thmm.bic, [12](#)  
thmm.bw.ctrl, [2](#), [10](#), [11](#), [13](#)  
thmm.estep, [2](#), [11](#), [14](#)  
thmm.forward, [9](#), [10](#), [12](#), [15](#), [16](#), [17](#), [22](#)  
thmm.forward.backward, [10](#), [14](#), [15](#), [16](#)  
thmm.init, [9](#), [10](#), [12–16](#), [17](#), [20–22](#)  
thmm.parameters, [19](#)  
thmm.simulate, [20](#)  
thmm.viterbi, [21](#)